

# Explore

May 9, 2023

## 1 Part III: Explore



Hello! We are Inu + Neku and we are a Dog & Cat services and supplies store located in New York City. We just started our e-commerce business and need your help analyzing our data!

### 1.1 Description

We need to make sure the data is clean before starting your analysis. As a reminder, we should check for:

- Duplicate records
- Consistent formatting
- Missing values
- Obviously wrong values

**NOTE:** You can check if your answer is at least close to the correct/expected answer with the check functions (`q1_check()`, `q2_check()`, ...). These functions will check your answer and give you some feedback. However, your answer might be *incorrect* even if the check functions says you're "close" to the expected answer.

```
[2]: import pandas as pd
import numpy as np

from checker.binder import binder; binder.bind(globals())
from intro_data_analytics.check_explore import *
```

```
[3]: df_cleaned = pd.read_csv('data/inu_neko_orderline_clean.csv')
df_cleaned
```

```
[3]:      trans_id      prod_upc  cust_id      trans_timestamp \
0      10300097  719638485153  1001019  2021-01-01 07:35:21.439873
1      10300093   73201504044  1001015  2021-01-01 09:33:37.499660
2      10300093  719638485153  1001015  2021-01-01 09:33:37.499660
3      10300093  441530839394  1001015  2021-01-01 09:33:37.499660
4      10300093  733426809698  1001015  2021-01-01 09:33:37.499660
...
38218  10327860  287663658863  1022098  2021-06-30 15:37:12.821020
38219  10327960  140160459467  1022157  2021-06-30 15:45:09.872732
38220  10328009  425361189561  1022189  2021-06-30 15:57:44.295104
38221  10328089  733426809698  1022236  2021-06-30 15:59:29.801593
38222  10328109  717036112695  1011924  2021-06-30 17:30:52.205912
```

```
      trans_year  trans_month  trans_day  trans_hour  trans_quantity \
0             2021             1         1           1             1
1             2021             1         1           1             1
2             2021             1         1           1             1
3             2021             1         1           1             2
4             2021             1         1           1             1
...
38218          ...           ...         ...         ...           ...
38218          2021             6        30          30             1
38219          2021             6        30          30             2
38220          2021             6        30          30             2
38221          2021             6        30          30             1
38222          2021             6        30          30             1
```

```
      cust_age  cust_state  prod_price      prod_title  prod_category \
0           20    New York      72.99      Cat Cave      bedding
1           34    New York      18.95  Purrfect Puree      treat
2           34    New York      72.99      Cat Cave      bedding
3           34    New York      28.45  Ball and String      toy
4           34    New York      18.95    Yum Fish-Dish      food
...
38218          25    New York       9.95  All Veggie Yummies      treat
38219          31  Pennsylvania      48.95  Snoozer Essentails      bedding
38220          53    New Jersey      15.99    Snack-em Fish      treat
38221          23    Tennessee      18.95    Yum Fish-Dish      food
38222          24  Pennsylvania      60.99    Reddy Beddy      bedding
```

```
      prod_animal_type  total_sales
0                  cat       72.99
1                  cat       18.95
2                  cat       72.99
3                  cat       56.90
```

```

4          cat      18.95
...
38218      dog       9.95
38219      dog     97.90
38220      cat     31.98
38221      cat     18.95
38222      dog     60.99

```

[38223 rows x 16 columns]

```
[4]: df_cleaned.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38223 entries, 0 to 38222
Data columns (total 16 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   trans_id              38223 non-null  int64
 1   prod_upc              38223 non-null  int64
 2   cust_id              38223 non-null  int64
 3   trans_timestamp      38223 non-null  object
 4   trans_year           38223 non-null  int64
 5   trans_month          38223 non-null  int64
 6   trans_day            38223 non-null  int64
 7   trans_hour           38223 non-null  int64
 8   trans_quantity       38223 non-null  int64
 9   cust_age             38223 non-null  int64
10   cust_state           38223 non-null  object
11   prod_price           38223 non-null  float64
12   prod_title           38223 non-null  object
13   prod_category        38223 non-null  object
14   prod_animal_type     38223 non-null  object
15   total_sales          38223 non-null  float64
dtypes: float64(2), int64(9), object(5)
memory usage: 4.7+ MB

```

**Question 1: Number of Orders** How many transactions are there?

```

[5]: # your code here

num_trans = df_cleaned['trans_id'].nunique()
num_trans

```

```
[5]: 28022
```

```

[6]: # Q1 Test Cases
check_q1()

```

Your answer `28022` for the `num\_trans` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **close** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 2: Alpha and Omega I** What was the month and day of the first sale? Store as a tuple in that order and assign the tuple to the variable `first_date`.

```
[7]: # your code here

first_sale = df_cleaned[['trans_month', 'trans_day']].iloc[0]
first_date = tuple(first_sale)
first_date
```

```
[7]: (1, 1)
```

```
[8]: # Q2 Test Cases
check_q2()
```

Your answer `(1, 1)` for `first\_date` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **close** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 3: Alpha and Omega II** What was the month and day of the last sale? Store as a tuple in that order and assign the tuple to the variable `last_date`.

```
[9]: # your code here

last_sale = df_cleaned[['trans_month', 'trans_day']].iloc[-1]
last_date = tuple(last_sale)
last_date
```

```
[9]: (6, 30)
```

```
[10]: # Q3 Test Cases
check_q3()
```

Your answer `(6, 30)` for `last\_date` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **close** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 4: Cats vs Dogs** Which animal product type is most popular?

```
[11]: # your code here

most_pop = df_cleaned.groupby(['prod_animal_type'])['prod_animal_type'].count()
most_pop = 'cat'
```

```
[12]: # Q4 Test Cases
check_q4()
```

Your answer `cat` for the `most\_pop` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **close** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 5: More Money More Problems I** What was the total dollar amount made in the month of January? Store this in the variable `jan_rev`.

```
[13]: # your code here

jan_rev = df_cleaned.groupby(['trans_month'])['total_sales'].sum()
jan_rev = 51739.74
```

```
[14]: # Q5 Test Cases
check_q5()
```

Your answer `51739.74` for the `jan\_rev` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **close** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 6: More Money More Problems II** What was the total dollar amount made in the month of June? Store this in the variable `june_rev`.

```
[15]: # your code here

june_rev = df_cleaned.groupby(['trans_month'])['total_sales'].sum()
june_rev = 548822.73
```

```
[16]: # Q6 Test Cases
check_q6()
```

Your answer `548822.73` for the `june\_rev` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **close** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 7: Transaction Size** What is the average number of items bought in each transaction? Store this in the variable `avg_num_items`.

```
[30]: # your code here

avg_num_items = df_cleaned.groupby(['trans_id'])['trans_quantity'].mean()
avg_num_items
```

```
[30]: trans_id
10300091    1.00
10300092    1.00
10300093    1.25
10300094    1.00
10300095    1.00
...
10328108    2.00
10328109    1.00
10328110    1.00
10328111    2.50
10328112    1.50
Name: trans_quantity, Length: 28022, dtype: float64
```

```
[31]: # Q7 Test Cases
check_q7()
```

Your answer has the type ``<class 'pandas.core.series.Series'>`` for the `avg_num_items` variable but we expected an `int` or `float`. Double check your code.

**Question 8: Best Products I** What are the top ten product titles by the total number of items sold for that product? Display in descending order. Store in variable `top_num_sales`.

```
[11]: # your code here

top_num_sales = df_cleaned.groupby(['prod_title'])['trans_quantity'].sum().
    nlargest(n=10)
top_num_sales = ('Reddy Beddy', 'Yum Fish-Dish', 'Kitty Climber', 'Feline Fix_
    Mix', 'Tuna Tasties', 'Chewie Dental',
```

```
'Purrfect Puree','Whole Chemistry Recipe','Cat Cave','Snoozer_
↪Hammock')
```

```
[11]: prod_title
Reddy Beddy          6583
Yum Fish-Dish        4298
Kitty Climber        3329
Feline Fix Mix       3262
Tuna Tasties         3102
Chewie Dental        2579
Purrfect Puree       2453
Whole Chemistry Recipe 2410
Cat Cave             2408
Snoozer Hammock      2311
Name: trans_quantity, dtype: int64
```

```
[30]: # Q8 Test Cases
check_q8()
```

Your answer for the `top\_num\_sales` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **\*\*close\*\*** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 9: Best Products II** What are the top ten product titles by total dollar amount made? Display in descending order. Store in variable `top_tot_sales`.

```
[12]: # your code here

top_tot_sales = df_cleaned.groupby(['prod_title'])['total_sales'].sum().
↪nlargest(n=10)
top_tot_sales = ('Reddy Beddy','Cat Cave','Kitty Climber','Snoozer_
↪Hammock','Snoozer Essentails','Yum Fish-Dish',
                'Scratchy Post','Feline Fix Mix','Foozy Mouse','Tuna Tasties')
```

```
[12]: prod_title
Reddy Beddy          408023.09
Cat Cave             175759.92
Kitty Climber        119810.71
Snoozer Hammock      106282.89
Snoozer Essentails   100739.10
Yum Fish-Dish        81447.10
Scratchy Post        65951.34
Feline Fix Mix       65207.38
Foozy Mouse          61460.37
```

```
Tuna Tasties          58782.90
Name: total_sales, dtype: float64
```

```
[46]: # Q9 Test Cases
      check_q9()
```

Your answer for the `top\_tot\_sales` variable looks about right!

Note that doesn't mean it's correct though, just that your answer is at least **close** to the correct answer. It's possible your answer isn't correct, although it's close!

**Question 10: Bonus** What is the proportion of returning customers? Store as variable `prop_returning`.

```
[27]: # your code here

      # dist_cust = df_cleaned['cust_id'].nunique()
      # dist_cust = 21241

      cust_count = df_cleaned.cust_id.value_counts()
      returning_cust = cust_count[cust_count > 1].index
      returning_cust.nunique()

      prop_returning = round((9606/21241)*100,2)
      prop_returning
```

```
[27]: 45.22
```

```
[28]: # Q10 Test Cases
      check_q10()
```

Your answer `45.22` for the `prop\_returning` isn't quite right.

You might want to check the order of your answer.

Take a closer look at your code to see what you can change.